

# When Does Adapting the Meta-Parameter Pay? Minimal Governors and Structural Limits for Weight- and Horizon-Governed MPC

Víctor Valdivieso<sup>a</sup>, Francisco G. Montoya<sup>a</sup>, Alejandro Campa-Pinto<sup>b</sup>, and Francisco Manuel Arrabal-Campos<sup>a,\*</sup>

**Abstract**—A slow outer loop that retunes an inner model predictive controller — its scalarization weight, or its prediction horizon — is an appealing idea, and we set out to make an elaborate version of it work: a second-order wave field on the interconnection graph, transporting each subsystem’s difficulty to its neighbours. It lost in every physically grounded arena we built. What replaces it is sharper than a negative result. First, two minimal governors that do pay. In a chain with a physically travelling disturbance, an *affine* delay line — perhaps delay minus a constant anticipation term — recovers 91% of an acausal clairvoyant reference — bootstrap 95% interval [84.9, 103.9], which contains 100% — and cuts closed-loop cost by 12.4% against the best possible local detector; the field, a shaped filter, and the proportional delay family that most designs default to are all significantly worse, and the missing constant term is precisely why. In a fleet of loops sharing one embedded controller under a hard computation budget, a governor that requests the maximum horizon *only* when the relaxed-dynamic-programming certificate breaks, and decays slowly otherwise, beats the best fixed horizon at equal *spent* computation by up to 6% — at every budget that binds when disturbances are independent, and at 4 of eight budgets under a perfectly correlated front (Holm-adjusted). A decomposition into its two mechanisms shows them to be complementary rather than cumulative: cross-loop allocation pays only when disturbances are independent, temporal adaptation only under a common front — where nothing can be reallocated but the moment to spend is well defined — and each is neutral or harmful in the other’s regime. Second, structural limits that explain why nothing more elaborate could win: amplitude-scale invariance of the optimal meta-parameter, the linear-time-invariance of the graph field, the DC cross-gain of its forcing, and a quantization-limited reach law. Third, the part most likely to be useful beyond our own problem: five experimental leaks, two acausal, two of accounting, one of channel, every one of which inverted a conclusion of ours before adversarial review caught it, together with the cheap executable tests that now keep each of them out.

**Index Terms**—Predictive control, adaptive horizon, resource-aware control, distributed control, experimental methodology, negative results.

## I. INTRODUCTION

<sup>a</sup>Department of Engineering, Research Centre CIAIMBITAL, University of Almería, 04120 Almería, Spain.

<sup>b</sup>Department of Computer Science and Electronics, Universidad de la Costa, Barranquilla 080002, Colombia.

\*Corresponding author (e-mail: fmarrabal@ual.es).

Every numerical claim below is regenerated by archived scripts. The defects found during four rounds of adversarial review of this work — five of them serious enough to invert a conclusion — are each encoded as a regression assertion in an executable audit suite distributed with the code: <https://github.com/fmarrabal/limits-horizon-governed-mpc>.

MODEL predictive control exposes a handful of design quantities that are usually chosen once and then frozen: how much one objective matters relative to another, how far ahead to predict, how hard to push soft constraints. Letting these *meta-parameters* vary online, driven by some measured indicator of how hard the current situation is, is a natural thing to want, and adaptive-horizon MPC, resource-aware MPC and priority scheduling all do a version of it. A distributed version suggests itself just as readily: couple the meta-parameters of neighbouring subsystems through a diffusion or wave operator on the interconnection graph, so that one subsystem’s decision reaches the others without anyone having to negotiate.

We set out to make that work, across six arenas built for the purpose. The elaborate governor won in exactly one of them — a synthetic bench whose priority signal had no physical origin, which is why we do not count it — and in none of the physically grounded ones. This paper reports what survived instead, and it turned out to be more useful than the idea we started with.

The first thing that survived is a pair of *minimal* governors that genuinely pay, each beating not only its elaborate rivals but the natural simple ones, and each for a reason that can be stated in one sentence. In a transport-coupled chain, upstream information is worth 12.4% of closed-loop cost, and the mechanism that captures it is an affine delay line: what matters is that every node receives its warning with the *same* anticipation, which a proportional delay cannot provide and a constant term can. In a fleet of loops sharing one embedded controller, governing the horizon is worth up to 7% at equal spent computation, and the governor that achieves it is simply a detector of certificate violation behind a slow leak: two effective parameters. In both cases everything more sophisticated, including the field we were defending, is statistically no better or significantly worse.

The second thing that survived is a set of structural statements that say, in advance, why the sophisticated alternatives could not have won (Section III). They are elementary: a homogeneity argument, the observation that a linear time-invariant operator is a fixed reshaping, a DC-gain computation. Each one, once stated with its hypotheses complete, retired an entire family of test benches or mechanisms, several of them ours.

The third is methodological, and we believe it travels furthest (Section VII). Every arena we built contained at least one defect that inverted or invalidated its first conclusion. Five of those defects form a pattern worth naming: an acausal time

shift hiding in a tuning grid, an acausal normalization hiding in a scale factor, a broken channel hiding under a plausible forcing term, and two accounting slips, one of which inflated the result it touched while the other, more insidiously, *deflated* it. Nearly all were found by automated adversarial review rather than by us; one we found ourselves, late. Each is now an executable assertion, and the tests that catch them (a pulse probe, a truncation probe, a DC-gain probe, spent-computation accounting, boundary-optimum rejection) are cheap enough to run inside a tuning loop.

The weight governor of Section II rests on a multiparametric fact whose sign is stated the other way round in the formulation we build on; we use the corrected version, and treat the correction itself — proof, consequences and an exact repair — in a separate Comment [1].

### Relation to existing work

Feedforward compensation of a *known* transport delay is classical, and it is standard practice in the domains our transport arena abstracts: district heating networks [2] and distributed solar collector fields [3], [4]. We do not claim novelty for the idea that delayed upstream information helps. The contribution is narrower and, we think, more transferable: when the plant model is kept nominal and the only channel available is the meta-parameter, the *structure* of the delay family decides everything — the proportional family that a per-hop design naturally produces is structurally unable to deliver constant anticipation, an affine family is, and nothing richer (second-order shaping, a wave field on the graph) improves on it. We found this because our own “unreachable oracle” turned out to be implementable, which is itself a methodological point.

Distributed and hierarchical MPC architectures are surveyed in [5], [6], with plug-and-play variants in [7] and canal networks in [8]. Abstracting a chain of subsystems by a PDE and reasoning about propagation on it appears in vehicular platooning [9], and wave-equation-based formation control exists [10]; our negative results on the graph field do not touch those uses, which couple the *physical* states rather than the meta-parameters. On horizon adaptation, variable and adaptive horizons with guarantees are treated in [11], [12], and a robust self-triggered scheme in which the horizon adapts as part of the triggering mechanism appears in [13]. We contribute no new guarantee; we characterize empirically *which* governor decodes the suboptimality signal, and we connect the winner to the certificate itself.

### Contributions

- 1) **The affine-delay result** (Section IV): in a transport-coupled storage chain, constant-anticipation delivery through the meta-parameter channel recovers 90.9% of an acausal clairvoyant reference, an interval wide enough to contain the reference itself; the proportional-delay family, a matched shaped filter, and a wave field are structurally or statistically inferior, with the failure of the proportional family explained.

- 2) **The certificate-detector result** (Section V): a horizon governor that maps certificate violation ( $\alpha_t \leq 0$ ) to maximal horizon and decays otherwise sits below the fixed-horizon frontier in a single loop and beats the best fixed horizon by up to 6% in a fleet at equal spent computation, under both disturbance-correlation regimes; with a decomposition, one arm per mechanism, showing that allocation and timing carry one regime each.
- 3) **Structural limits** (Section III) with hypotheses stated in full and their necessity demonstrated by counterexample.
- 4) **An audit methodology** (Section VII): five real leaks, the conclusion each inverted, and the executable test that now guards each; all distributed as a regression suite whose every assertion encodes a defect found at least once.

## II. PROBLEM SETTING

Consider the discrete-time linear system

$$x_{k+1} = Ax_k + Bu_k + w_k, \quad x_k \in \mathbb{R}^n, \quad u_k \in \mathcal{U}, \quad (1)$$

with  $\mathcal{U} = \{u : \|u\|_\infty \leq u_{\max}\}$  and  $w_k$  an unmeasured disturbance. Throughout, the MPC is *nominal*: it predicts with  $x^+ = Ax + Bu$  and no disturbance model, which is both standard and the honest setting for our question — with a disturbance model the controller would compensate by feedforward and there would be nothing left for a meta-parameter to arbitrate.

### A. Weight-governed multiobjective MPC

Two stage costs  $\ell_i(x, u) = x^\top Q_i x + u^\top R_i u$  compete, scalarized on the simplex with weight  $(1 - \alpha, \alpha)$ ,  $\alpha \in [0, 1]$ . With terminal weights  $P_i$  from the discrete Lyapunov equation for a common terminal gain  $K_f$  and an invariant terminal set from the maximal output admissible set construction [14], the controller solves

$$V^*(x, \alpha) = \min_{U \in \mathcal{U}^{N \cap \mathcal{X}_N(x)}} (1 - \alpha)J_1(x, U) + \alpha J_2(x, U), \quad (2)$$

and the applied weight is restricted by the contractive constraint inherited from the weighted-sum multiobjective formulation of [15],  $V^*(x, \alpha) \leq J_a$  with  $J_a$  the cost of the shifted sequence under the previous weight. One parametric fact matters for implementation: for fixed  $x$  the map  $\alpha \mapsto V^*(x, \alpha)$  is a pointwise minimum of functions affine in  $\alpha$ , hence *concave* [16, Sec. 5.6.1, Ex. 3.9], so the admissible weight set  $\mathcal{A}(x, J_a) = \{\alpha : V^*(x, \alpha) \leq J_a\}$  is in general non-convex — in our instances it was non-convex at roughly one fifth of the sampled contraction levels — and must be enumerated piece by piece (one linear program per active piece, the exact repair of [1]) or, in the scalar case used here, gridded, rather than projected onto. That sign is the opposite of the one stated in [15, Lemma 4]; we use the corrected version, and the correction itself — its proof, its consequences for Theorems 6 and 7 of that formulation, and an exact repair at the same computational cost — is the subject of a separate Comment [1].

### B. Horizon-governed MPC and the suboptimality index

For the single-objective case we use the relaxed dynamic programming index of Grüne and Pannek [17]–[19]. With  $V_N$  the optimal value of the  $N$ -step problem *without* terminal ingredients and  $\ell_t$  the incurred stage cost,

$$\alpha_t = \frac{V_N(x_t) - V_N(x_{t+1})}{\ell_t}, \quad (3)$$

and a *uniform*  $\alpha_t \geq \alpha > 0$  certifies  $J_\infty \leq V_N/\alpha$  [18]. Evaluated along a closed loop, (3) is an a posteriori diagnostic, not a per-step certificate: it can be negative, and the telescoping identity it satisfies is accounting rather than a guarantee. The event  $\alpha_t \leq 0$  — the certificate breaking — turns out to be the one feature of this signal a governor should react to (Section V).

Because (3) is a ratio it is uninformative when  $\ell_t \approx 0$ , and gating on  $\ell_t$  alone silently annuls exactly the sample created by a disturbance kick, whose pre-kick stage cost is essentially zero; our first implementation did precisely that and produced governor traces such as  $N = [2, 2, 16, 2, 16, \dots]$ . A jump in  $V_N$  betrays an event even when  $\ell_t$  vanishes, so the informativeness gate must be bilateral.

### C. The graph field, i.e. the elaborate governor under test

The distributed governor we set out to defend assigns each subsystem  $i$  a field value  $h_i$  that becomes its meta-parameter request. With  $L$  the graph Laplacian and  $A$  the adjacency matrix,

$$K = \omega_0^2 I + c^2 L, \quad C = 2\zeta\omega_0 I + DL, \quad G = bA + \beta A^{\circ 3}, \quad (4)$$

a wave branch  $\ddot{h} = -(C + G)\dot{h} - Kh + f(h_d)$  integrated by Störmer–Verlet, and a diffusion branch  $\gamma\dot{h} + (K + G)h = f(h_d)$  integrated by an IMEX scheme. The two branches share  $K$ , hence the steady-state profile, so wave-versus-diffusion is a mechanism null by construction. The choice of forcing  $f(\cdot)$  looks like a detail; Proposition 3 shows it decides whether the field can transport anything sustained at all.

## III. STRUCTURAL LIMITS

The statements in this section are elementary. We give them anyway, with hypotheses complete, because each one retired a family of benches or mechanisms, and because in two cases the hypothesis we originally omitted turned out to be the load-bearing one.

### A. Amplitude-scale invariance

**Proposition 1.** Consider (1) with quadratic stage costs, zero initial state  $x_0 = 0$ , and every constraint — input, state, terminal — inactive along the whole trajectory, for every meta-parameter value swept and every scale considered. Fix the meta-parameter (weight  $\alpha$  or horizon  $N$ ). If the disturbance is scaled  $w \mapsto \lambda w$ ,  $\lambda > 0$ , the optimal trajectories scale by  $\lambda$ , the cost by  $\lambda^2$ , and therefore the entire normalized cost curve  $J(\cdot)/\lambda^2$  — hence also its set of minimizers — is independent of  $\lambda$ .

*Proof.* With no active constraint the closed loop is linear in the exogenous input for each fixed meta-parameter, and with  $x_0 = 0$  there is no free response to break homogeneity; a quadratic cost then scales by  $\lambda^2$ , and a positive common factor moves nothing.  $\square$

Numerically, with both optima interior to their grids ( $\alpha^* = 0.20$ ,  $N^* = 32$ ), the normalized cost curves for  $\lambda \in \{0.5, 1, 2, 4, 8\}$  coincide to relative dispersion  $5.4 \times 10^{-14}$  (weight) and  $1.2 \times 10^{-16}$  (horizon) (Fig. 1). The hypotheses are not decorative: with  $x_0 = (0, -4)$  on the same strictly linear plant — one point, not a family, and reported as such — the normalized cost curves stop coinciding and disperse by 0.999, order one against the  $10^{-14}$  of the hypothesis-satisfying case. The argmin, however, did *not* move at that point ( $N^* = 8$  at both scales), and the archived assertion records exactly that. We keep the counterexample because the proposition is about the curve; a reader who tested it on argmins alone would have concluded the hypotheses were decorative. And the statement is about the *curve*, deliberately: when the minimum is flat, the literal argmin can hop between points whose costs differ by  $10^{-16}$ , so comparing argmins is the wrong test — our own audit assertion failed for exactly that reason until we rewrote it to compare curves and near-optimal sets.

What follows is bounded, and our first draft over-claimed it. Homogeneity forbids the optimal meta-parameter from depending on the *absolute amplitude* of the disturbance; it does not forbid dependence on the shape or phase of a normalized envelope, which are scale-invariant. So a bench is empty exactly when its demand signal is *postulated as a function of absolute amplitude* and its evaluation metric is homogeneous of degree two. Our narrow-band arena satisfies the first condition and *not* the second: it is evaluated with a superlinear fatigue cost  $q e^2 + r|u|^4$ , which is not homogeneous of degree two, so Proposition 1 does not apply to it — and could not, since the schedule  $\alpha_d = 1/(1 + \kappa E^2)$  tested there is *derived* from precisely that non-homogeneity. That arena lost in all eight configurations (Fig. 6) for a different and more mundane reason, which we state because a defeat of 41 to 134% deserves a cause rather than a citation: the rival constant weight is re-optimized for each operating condition, so it already sits wherever the amplitude dependence would place it, and what the schedule adds on top is modulation *within* a condition, which the loop pays for and does not recover. The proposition and the arena point the same way, but only the arena is evidence here.

### B. The graph field is an LTI filter bank, with a scope warning

**Proposition 2.** With  $b = \beta = 0$  in (4), the map from the injected signal  $h_d$  to the node values  $h$ , in either branch, is linear and time invariant: each node’s response is the convolution of the source with a fixed impulse response. (Verified to  $1.1 \times 10^{-15}$  by superposition, exactly by time invariance, and to  $4.4 \times 10^{-16}$  by convolution reconstruction.)

Whatever the field contributes over the signal entering at the source is therefore a fixed linear reshaping — delay plus spreading. Two honest caveats, both of which we got wrong




Fig. 1. Amplitude-scale invariance and its limit. (a,b) Normalized cost  $J/\lambda^2$  against weight and horizon for five disturbance scales, both optima interior ( $\alpha^* = 0.20$ ,  $N^* = 32$ ): the curves coincide. (c) With a saturating actuator a level-dependent optimum does appear, but only in a narrow window of the disturbance envelope.

before getting them right. First, the field belongs to the class of LTI banks of order  $2M$ ; a rival parametrizing that whole class would contain it, but the seven-parameter shaped delay we use in Section IV does *not*: measured on interior hops of an eight-node chain, the field’s kernels broaden with hop count, by a factor that depends on the width measure used, at the operating point as well as at nominal parameters. The field’s defeat there is an empirical finding, not a corollary. Second, the deployed governor is *not* LTI — between  $h$  and the applied meta-parameter sit a saturation, a projection onto  $\mathcal{A}$ , and a quantizer — and the reach law below depends on the quantizer.

### C. The forcing decides whether anything crosses

**Proposition 3.** *Take the forcing  $f(h_d) = Kh_d$ , so the stiffness term reads  $-K(h - h_d)$  and each node relaxes to its own target. Then the equilibrium is  $h = h_d$  exactly,  $H(0) = K^{-1}K = I$ , and the cross-node DC gain is identically zero: a sustained alert injected at one node reaches no other node, under any tuning. With  $f(h_d) = \omega_0^2 h_d$  the equilibrium is  $\omega_0^2 K^{-1} h_d$ , whose off-diagonal entries are nonzero, and with  $G = 0$  both branches share the same  $H(0)$ , preserving the mechanism null.*

A unit step at node 1 of a six-node chain settles to  $[1, 0, 0, 0, 0, 0]$  to machine precision under the first forcing (Fig. 2); only edges cross, as transients. Our transport experiment originally asked the field to deliver a sustained window under exactly this forcing — the one thing its structure forbids — and the field’s tuned gains pinned themselves to the top of every grid trying to compensate, which was the observable symptom.

**Observation 1** (Quantization-limited reach). If the closed-loop transported amplitude decays geometrically with hop count,  $a_n = a_0 \rho^n$ , and the meta-parameter moves only when that amplitude exceeds one grid step  $\Delta$ , the reach is  $n_{\max} = \lfloor \log \Delta / \log \rho \rfloor$ : logarithmic in the resolution of the meta-parameter.

Measured on an eight-node chain with  $\rho = 0.358$  fitted on interior nodes, the step law matches at three of four resolutions and overpredicts by one hop at  $\Delta = 0.005$  — the third of the four, not the finest, which it gets right. At that resolution the geometric law puts the disputed hop’s amplitude at 0.0059 against a step of 0.005: a fifth of a step above threshold, so a single quantization event decides it (Fig. 3). We report the pre-registered verdict as the script recorded it — the law is *not* validated, four resolutions being too few to call one miss a coincidence. Two cautions we needed: the reach is an integer, so plotting the continuous expression against integer measurements makes a correct law look wrong; and the terminal node of a chain has degree one and retains signal, so fitting  $\rho$  through it biases the estimate — an artifact that, we admit, also contaminated one of our own intermediate re-measurements before we excluded terminal nodes throughout. The half-step threshold that round-to-nearest quantization suggests fits none of the four resolutions; the amplitude we report is a dispersion along the trace, not a peak, and the empirical criterion is the full step.

## IV. RESULT I: CONSTANT ANTICIPATION, NOT PROPORTIONAL DELAY

The one use of a graph field that the structural limits leave open is transport: a disturbance that travels physically downstream, so that the upstream neighbour holds information about the downstream node’s future which no local horizon can obtain. Node  $i$  experiences the source disturbance delayed by  $i\tau$  steps ( $\tau = 6$  in our chain of four nodes).

### A. Three viability filters

Three checks preceded any comparison, each able to kill the arena. Without a storage state, the lead sweep is flat: a nominal regulator at rest applies  $u^* = 0$  for every meta-parameter value, so advance warning has no channel to act through — unless the governor itself has lag, which is why we state this as a filter rather than a theorem. With storage — two objectives at different setpoints, so the weight selects the buffer level and




Fig. 2. Proposition 3. (a) Under the homeostatic forcing the DC transfer is the identity: a sustained input at the source produces nothing at any other node. (b) Under the source forcing the sustained component crosses, identically in both branches when  $G = 0$ .




Fig. 3. Observation 1. (a) Transported amplitude against hop count for four grid resolutions; crosses mark nodes quantized to exactly zero, dotted lines the grid step. (b) Measured reach against grid step, with the step law and  $\rho = 0.358$  fitted on interior nodes; the open marker is the one resolution where the law overpredicts by a hop.

anticipating means charging early — anticipation buys +4.0% over the ideal reactive limit ( $t = +5.80$ ,  $p = 2.6 \times 10^{-4}$ , nine of ten seeds, tuning and evaluation seeds disjoint). The useful lead is short, two to four steps across the physical range, and grows with actuator sluggishness as the mechanism predicts.

### B. The discovery: our unreachable oracle was implementable

We compared delivery mechanisms for the source measurement: an ideal local detector (learns of the event the instant it arrives — stronger than any realizable detector), a per-hop delay line  $k_i = \delta i$ , the wave field, a shaped delay (per-hop delay plus a second-order stage, matched to the field’s seven parameters), and, as a reference ceiling, a “clairvoyant oracle” that opens a gate a fixed lead before each arrival.

Getting this comparison honest took three corrections, described in Section VII; here is the one that produced the result. After the causality guard was in place, the tuned oracle — supposedly an unreachable bound — opened its gates a fixed

number of samples ahead of each arrival at every node, and never before the source measured. It was causal, and it was reproduced *exactly* (difference 0.0) by a delay line

$$k_i = \max\{0, \tau i - \ell\}, \quad (5)$$

affine in the hop index with a constant anticipation term  $\ell$ . The competition had excluded the best member of its own family. The proportional family  $k_i = \delta i$  cannot express (5): with  $\delta = \tau$  the anticipation is zero everywhere, and with  $\delta < \tau$  it grows with distance, so far nodes are warned too early and near nodes too late. What the physics supplies is anticipation *growing* with distance ( $i\tau$ ); what the loop needs is anticipation *constant* across nodes (saturated, for near nodes, at the maximum physically available); the constant term is the conversion between the two.

TABLE I

TRANSPORT CHAIN, 30 EVALUATION SEEDS. “MDE” IS THE MINIMUM DETECTABLE EFFECT AT POWER 0.8 FOR THE PAIRED CONTRAST; CONTRASTS ARE HOLM-CORRECTED [20]. THE AFFINE DELAY AGAINST THE PROPORTIONAL FAMILY’S *best known* TUNING (FROM A DENSER GRID, ARCHIVED) STILL GIVES  $t = +5.40$ , SO THE DECISIVE CONTRAST DOES NOT DEPEND ON THE RIVAL’S VERSION.

| mechanism                    | cost (s.e.m.)        | vs. local    | % ceiling   |
|------------------------------|----------------------|--------------|-------------|
| acausal ceiling              | 48 944 (2828)        | +13.7        | 100.0       |
| affine + shaping (7 par.)    | 49 455 (2175)        | +12.8        | 93.4        |
| <b>affine delay (4 par.)</b> | <b>49 650 (2370)</b> | <b>+12.4</b> | <b>90.9</b> |
| wave field (7 par.)          | 51 054 (2083)        | +9.9         | 72.8        |
| proportional delay (3 par.)  | 51 711 (2479)        | +8.8         | 64.3        |
| ideal local detection        | 56 690 (1301)        | 0.0          | 0.0         |

*paired contrasts, Holm;  $\delta$ ,  $t$ ,  $p$ , MDE as % of cost*

|                         |                                         |
|-------------------------|-----------------------------------------|
| affine vs. proportional | +2060, $t=+17.87$ , $p<10^{-4}$ *, 0.6% |
| affine vs. wave field   | +1403, $t=+4.64$ , $p=0.0002$ *, 1.5%   |
| shaping the affine      | −195, $t=−0.97$ , $p=0.34$ , 1.0%       |
| affine vs. local        | +7040, $t=+6.29$ , $p<10^{-4}$ *, 5.5%  |
| ceiling vs. affine      | +706, $t=+1.49$ , $p=0.29$ , 2.3%       |

### C. The result

Table I and Fig. 4 give the final comparison: thirty evaluation seeds disjoint from the three tuning seeds, causality guard active inside the tuning loop, causal (parameter-only) normalization, equal tuning-grid densities for all mechanisms, and every tuned optimum interior to its grid with one declared exception: the proportional family’s high weight lands on  $a_{hi} = 1$ , the top of the simplex, which is a physical limit and not a grid artifact — and it is the *rival* whose optimum sits there, so the exception does not favour our mechanism. The archived sweep records the boundary status of every tuning, and the build check re-reads it. The acausal reference itself sits at both simplex limits ( $a_{lo} = 0$ ,  $a_{hi} = 1$ ) by construction: it is a reference, not a tuned mechanism, and the rule does not apply to it. The genuinely acausal ceiling uses per-node anticipation beyond what the source can supply and is labelled as a reference, not a mechanism and not a bound: a causal mechanism turned out to reach it, which is the point of Section IV.

Three readings. Having upstream information is worth 12.4% against the best possible local detector, and the study can resolve effects an order of magnitude smaller than that. The *structure* of delivery decides: affine beats proportional at  $t = +17.87$  under equal tuning densities, and at  $t = +5.40$  even against the proportional family’s best tuning from a denser grid. And nothing richer helps: adding the second-order shaping stage changes nothing resolvable ( $t = −0.97$  against an MDE of 1.0%), the wave field is significantly worse than the plain affine line ( $t = +4.64$ ), and the affine line is within  $−0.3\%$  to  $+3.7\%$  of the acausal reference (95% interval; it is cheaper than the reference in 13 of 30 seeds, paired  $t = +1.49$ , raw  $p = 0.15$ , Holm  $p = 0.29$ ; MDE 2.3% — an equivalence bound, not a proof of optimality). Resampling the thirty paired seeds puts the recovered fraction at 90.9% with a 95% interval of [84.9, 103.9], and in 6.3% of the resamples the causal line is *cheaper* than the acausal reference. That is why we report a reference rather than a bound: a quantity a causal mechanism

TABLE II

HORIZON GOVERNANCE, OUT-OF-ENVELOPE SCENARIO, TEN SEEDS. COMPUTATION, COST AND ASSIGNMENT CORRELATION ARE MEANS OVER THE SEEDS THAT DID *not* DIVERGE; THE LAST COLUMN COUNTS THE DIVERGENCES. DIVERGENT RUNS ARE ALSO EXCLUDED PAIRWISE FROM THE CONTRASTS; THE  $N = 2$  ROW’S COMPUTATION AND COST ARE MEANS OVER A RUN SET WITH HALF THE SEEDS DEAD, AND ARE SHOWN FOR COMPLETENESS RATHER THAN COMPARISON.

| governor                        | comput.     | cost         | corr  | div. |
|---------------------------------|-------------|--------------|-------|------|
| fixed $N = 2$                   | 610         | 29653        | —     | 5    |
| fixed $N = 4$                   | 1600        | 12561        | —     | 0    |
| fixed $N = 6$                   | 2400        | 12124        | —     | 0    |
| fixed $N = 8$                   | 3200        | 12104        | —     | 0    |
| fixed $N = 16$                  | 6400        | 12103        | —     | 0    |
| <b>leaky integrator</b>         | <b>1775</b> | <b>12211</b> | 0.935 | 0    |
| threshold                       | 3051        | 12103        | 0.858 | 0    |
| 2nd-order field ( $\zeta=0.5$ ) | 3133        | 17862        | 0.270 | 2    |
| memoryless map                  | 3666        | 27620        | 0.367 | 5    |

can beat six times in a hundred is not bounding anything. We state the one asymmetry that survives: the field’s kernel does differ across hops (Proposition 2 discussion), so its defeat is empirical, bounded by these MDEs, not structural.

## V. RESULT II: A CERTIFICATE DETECTOR GOVERNING THE HORIZON

### A. Single loop: the frontier, and who sits below it

We first close the loop  $\alpha_t \rightarrow \text{governor} \rightarrow N_t$  on a single plant, with pre-registered hypotheses, ten seeds, and computation accounted honestly: every QP is charged, including those used to measure  $\alpha_t$ , with the instrument solve reused as the next control solve when the horizon does not change. Table II reports the out-of-envelope scenario (larger kicks plus model mismatch) against the frontier traced by fixed horizons at equal computation.

Interpolating the frontier at the leaky integrator’s 1775 units gives roughly 12 470; the governor achieves 12 211, about 2% below the frontier, without knowing the correct horizon in advance. The second-order field rings — it drops the horizon mid-recovery, breaks the certificate again and re-excites itself; critical damping repairs its assignment correlation in the large-kick scenario ( $−0.42$  to  $+0.43$ ,  $t = +7.26$ ), which identifies ringing as the mechanism, but a critically damped second-order filter is a first-order filter with extra lag and does not beat one. On assignment correlation the leaky integrator is best in two of the four scenarios, including both out-of-envelope kick scenarios (with and without model mismatch). What broke out-of-envelope was not short *fixed* horizons at moderate  $N$  — replanning compensates from  $N = 4$  up, though  $N = 2$  does diverge on half the seeds — but short-*memory adaptive* governors, which under-assign right after a kick. The risk is aggressive downward modulation, not a modest fixed horizon.

The demand signal  $\alpha_t$  generates a pulse train; the natural decoder of pulse durations is a leaky integrator, not a resonator. Pushing the governor’s reference toward zero turned out to improve it monotonically, with a plateau at the limit: the optimal governor is a *detector of certificate violation* — request  $N_{\max}$  when  $\alpha_t \leq 0$ , request  $N_{\min}$  otherwise — behind






Fig. 4. (a) Anticipation received by each node under the two delay families: the proportional family either wastes the near nodes’ information or floods the far nodes early; the affine family delivers constant anticipation, saturated at what is physically available. (b) Final comparison, thirty seeds, error bars one s.e.m.: the affine delay is within  $-0.3\%$  to  $+3.7\%$  (95% interval) of the acausal reference, and neither shaping nor the wave field improves on it.

a slow leak ( $\tau = 30$  samples in the fleet; the single-loop bank of Table II ran with  $\tau = 14.7$  and  $a_{\text{ref}} = 0.6$ , and is therefore not this detector limit). Two effective parameters, and the event it reacts to is the one event in the signal with a meaning in the theory [17].

### B. A fleet sharing one embedded controller

A factor-two saving in computation matters most where predictive control is most constrained: one embedded controller serving many loops, the sampling period fixing a hard computation budget per step. We model  $M = 8$  thermal collector loops (fluid state coupled to a slower absorber,  $\tau_f = 2$  min,  $\tau_m = 5$  min,  $dt = 30$  s), with cloud-passage disturbances and a knob  $\rho$  setting the fraction that sweeps all loops at once. The penalized variable is the *absorber* temperature, which the flow reaches only through the fluid; without that delay between action and effect the bench is empty — penalizing the directly actuated state leaves the fixed-horizon frontier flat in  $N$  and nothing to govern, against a ratio  $J(N=2)/J(N=16) = 1.47$  here.

Two accounting disciplines decide whether this comparison means anything, and we broke both before fixing them (Section VII). The budget must bind on computation actually *spent* — including the certificate measurement — not on horizons granted; and the governor must be *tuned* under the same accounting it is evaluated under. With both in place, the rival is the fixed-horizon frontier interpolated, per seed, at the governor’s real computation; interpolation between adjacent integer horizons is achievable by alternating horizons between steps, so it concedes nothing.

Table III and Fig. 5: with independent disturbances the governor beats the best fixed horizon at equal spent computation at every budget that binds — seven of the eight rows, of which the three loosest are one experiment, since

TABLE III  
FLEET OF  $M = 8$  LOOPS, EIGHT EVALUATION SEEDS DISJOINT FROM THE THREE TUNING SEEDS. GOVERNOR TUNED BY A GRID SWEEP ARCHIVED WITH THE SCRIPT, UNDER THE EVALUATION METRIC ITSELF ( $\tau = 30$ , INTERIOR TO THE GRID; REFERENCE  $a_{\text{ref}} \rightarrow 0$ , THE CERTIFICATE-DETECTOR LIMIT, ON A PLATEAU WHOSE DISPERSION BELOW 0.03 IS 0.65%). COST AGAINST THE FIXED-HORIZON FRONTIER AT EQUAL SPENT COMPUTATION, PAIRED PER SEED AGAINST THAT SEED’S OWN FRONTIER. THE  $B = 16$  ROW IS DEGENERATE: EVERY LOOP IS PINNED AT  $N_{\min}$ .

| $B/(MN_{\max})$                                | spent | governed | frontier | change  |
|------------------------------------------------|-------|----------|----------|---------|
| <i>independent disturbances</i> ( $\rho = 0$ ) |       |          |          |         |
| 0.19                                           | 26.0  | 2814     | 2993     | −6.0% * |
| 0.25                                           | 35.2  | 2567     | 2732     | −6.0% * |
| 0.31                                           | 41.1  | 2470     | 2598     | −4.9% * |
| 0.38                                           | 42.9  | 2443     | 2568     | −4.9% * |
| $\geq 0.50$                                    | 43.1  | 2440     | 2564     | −4.8% * |
| <i>common front</i> ( $\rho = 1$ )             |       |          |          |         |
| 0.19                                           | 23.8  | 4589     | 4599     | −0.2%   |
| 0.25                                           | 29.1  | 4208     | 4343     | −3.1% * |
| 0.31                                           | 33.5  | 3912     | 4149     | −5.7% * |
| 0.38                                           | 37.3  | 3750     | 4003     | −6.3% * |
| 0.50                                           | 43.1  | 3645     | 3829     | −4.8% * |
| 0.75                                           | 48.2  | 3638     | 3716     | −2.1%   |
| 1.00                                           | 48.3  | 3638     | 3715     | −2.1%   |

the governor saturates at the same spent computation — by up to 6.0% (Holm-adjusted  $p < 10^{-4}$  throughout); under a perfectly correlated front it wins at 4 of eight budgets after Holm adjustment, by up to 6.3%, and the two loosest budgets favour it by 2.1% at raw  $p < 0.03$  but not after adjustment ( $p = 0.073$ ). Under independence the advantage is largest where the budget binds and settles at 4.8% once it stops binding; under a common front it is concentrated in mid-range budgets and vanishes at both ends. The eighth budget,  $B = 16$  (0.125), is omitted from the table: every loop is pinned at  $N_{\min}$

TABLE IV  
THE SIX ARENAS. “DEMAND” IS THE SIGNAL THE GOVERNOR IS ASKED TO TRACK.

| arena               | demand                         | outcome                                                         |
|---------------------|--------------------------------|-----------------------------------------------------------------|
| synthetic weights   | exogenous priority, no physics | field wins out of distribution; no physical origin, not counted |
| greenhouse price    | smooth, clean                  | modulation loses; the horizon already arbitrates the price      |
| graph, shared price | global signal                  | no gain: sharing what is already common                         |
| $\alpha_N$          | pulse train                    | pays; certificate detector + leak wins, the field rings         |
| horizon             | postulated                     | loses 8/8 to a per-condition re-optimized constant; quartic     |
| narrow-band         | from absolute                  | metric, outside Proposition 1                                   |
| vibration           | amplitude                      | pays 12.4%; affine delay                                        |
| transport, storage  | travelling event               | wins, field and shaping add nothing                             |

and the two arms coincide to 0.03%.

The decomposition in Fig. 5(b) separates the two mechanisms with one arm each, and it corrected our own reading of them. *Allocation* is isolated by keeping the governor and its temporal adaptation — the same total horizon per step — and only distributing it equally instead of in proportion to demand; *timing* is isolated by comparing that equal-split governed arm against a fixed horizon at equal spent computation. The two components turn out to be complementary rather than cumulative. Allocation pays +6.1% to +7.1% when disturbances are independent (all six budgets at equal spent computation, eight seeds each) and next to nothing under a common front (0.0% to +1.4%, significant after Holm at 2 of six budgets): with every loop in trouble at once there is no one to take horizon from. Timing does the opposite — it pays +2.2% to +5.2% under a common front at the 4 middle budgets, where the moment to spend is well defined, and nothing at the tightest and loosest; under independence it is not merely neutral but *harmful*, down to −3.7%: with someone always in trouble the aggregate demand is nearly constant, so modulating it in time only mistimes the spending that the allocation arm then has to repair. Each regime is carried by one mechanism, which is why the governor wins in both.

## VI. THE ARENA MAP

Table IV summarizes the six arenas. The one the field won is the one whose demand was an exogenous priority signal with no physical origin, and we do not count it as evidence.

There is a tension worth stating plainly, because it bounds future attempts. For modulation to pay one must cross a nonlinearity; crossing one makes the demand eventual, hence broadband, which is the regime where a leak beats a resonator. For inertia to pay one wants narrow-band demand, whose canonical manufacture is amplitude modulation, which Proposition 1 strips of absolute-amplitude content. The two conditions an inertial governor needs sit awkwardly together, and in six arenas they never met.

## VII. FIVE LEAKS, AND THE TESTS THAT NOW GUARD THEM

Every arena above was first decided by a defect rather than by its mechanisms. We document the five that form a pattern. None of them is exotic; each is the kind of line a careful person writes without noticing, and each has a mechanical guard cheap enough to run inside a tuning loop. Table V lists the resulting suite.

*a) Leak 1: a broken channel (DC cross-gain):* The field was first run with the homeostatic forcing, whose cross-node DC gain is identically zero (Proposition 3); the experiment asked it to deliver a sustained window, the one signal it cannot carry, and the field lost to a trivial delay line. The symptom was tuned gains pinned to grid boundaries. *Guard:* probe the channel with a sustained step and measure what arrives, before comparing anything (assertion A18).

*b) Leak 2: an acausal time shift:* A free time-shift in the signal-conditioning stage is a natural knob to include in a tuning grid; if the grid contains negative values, it silently hands the mechanism information from the future. Both leading mechanisms took the option — their gates opened up to eight steps before the source measured the event — and the only fair player was the simplest rival. The headline contrast of that version ( $t = +5.45$ ) dissolved to  $t = +0.05$  under the guard. *Guard:* a single-pulse probe inside the tuning loop; reject any configuration whose output moves before the input does.

*c) Leak 3: an acausal normalization:* The map from field value to weight was normalized by the trace’s min and max — computed over the *whole* trace, so the scale at time  $t$  depended on the future. Truncating the trace after the event measurably changed the gate *before* the event. Only the two seven-parameter mechanisms used the normalization; the delay line did not. *Guard:* the truncation probe — run the mechanism on the full and on a truncated input; on the shared prefix the outputs must agree exactly. This test requires no hypothesis about where the leak is, which is how it caught one we did not suspect. The fix (normalizing by the DC gain, a pure function of parameters) also removed a free knob.

*d) Leak 4: budget granted versus computation spent:* The fleet’s budget bound the horizons *granted*, while the governor also spent computation measuring its certificate; it consumed 18–26% more than the fixed rival at the same nominal budget, and “won” by overspending. *Guard:* account computation where it is spent and compare on the (spent, cost) plane. The single-loop experiment already did this; the fleet inherited the budget from the wrong side of the ledger.

*e) Leak 5: tuning under one metric, evaluating under another:* After fixing Leak 4 we re-tuned the governor by raw cost at fixed nominal budget — a criterion that still rewards overspending — and it moved the optimum to a faster leak and shrank the advantage at every budget that did not bind. Tuned under the evaluation metric itself (the gap to the per-seed frontier at spent computation, summed over the tuning budgets), the optimum returns to the slow leak and the governor wins by up to 6%. This leak is the mirror image of the others: it *deflated* the result it touched, and it did so twice — the tuning script of an earlier revision still used raw cost






Fig. 5. (a) Fleet cost against the fixed-horizon frontier at equal *spent* computation, for independent and perfectly correlated disturbances: the certificate-detector governor wins across budgets in both. (b) Decomposition into its two mechanisms, each isolated by one arm: allocation (demand-proportional versus equal-split, same governor and same total horizon per step) and timing (equal-split governed versus fixed horizon, at equal spent computation). They are complementary: allocation pays only under independence, timing only under a common front — up to +5.2% — and timing is actively harmful under independence.




Fig. 6. The narrow-band arena, built expressly to favour the field: resonant structure, beat disturbance, fatigue cost, and a schedule  $\alpha_d = 1/(1 + \kappa E^2)$  derived from that cost. It loses in all eight configurations against a per-condition re-optimized constant weight, which already absorbs the amplitude dependence the schedule was built to supply. The fatigue metric is quartic, so this arena lies outside Proposition 1: the defeat is empirical, not a corollary.

while the text claimed otherwise, which an adversarial review of the script, not of the text, caught. Leaks do not consistently favour the pet mechanism; they favour whatever the criterion mismatch happens to favour, which is why the guard must be mechanical. *Guard*: the tuning objective must be the evaluation metric, computed the same way.

f) *Three disciplines, used throughout*: Reject any tuned optimum on the boundary of its grid, with the exception of physical limits (a weight saturating at the simplex boundary is not a grid artifact) — applied to ourselves twice, including once to the evidence for Proposition 1. Report, with every null contrast, the minimum detectable effect: “no difference” in Table I means “bounded by the stated MDE,” never “proved zero.” And treat an “unreachable” bound that a guard fails to

TABLE V  
EXECUTABLE AUDIT SUITE: 21 CHECKS OVER 19 ASSERTIONS (A9 RUNS AT THREE DAMPING RATIOS). EACH ENCODES A DEFECT FOUND AT LEAST ONCE.

| id      | assertion                                                                                        |
|---------|--------------------------------------------------------------------------------------------------|
| A1–A3   | terminal inequality, invariance of $\Omega$ , constraint satisfaction                            |
| A4–A5   | $V^*(x, \cdot)$ concave in $\alpha$ ; admissible set can be non-convex                           |
| A6      | the certified radius is valid                                                                    |
| A7      | first/second-order matching is exact <i>in discrete time</i>                                     |
| A8      | anti-windup does not fire when nothing is blocked                                                |
| A9      | executed overshoot equals that of the <i>discrete</i> system                                     |
| A10–A11 | gyroscopic collocation stable; Merkin threshold                                                  |
| A12     | the inherited terminal ingredient fails (documented regression)                                  |
| A13–A14 | monotonicity duality of $V_N$ in $N$                                                             |
| A15     | scale invariance: normalized cost <i>curves</i> coincide; fails on boundary optima (Prop. 1)     |
| A16     | anticipation unusable without storage (viability filter)                                         |
| A17     | the field with $\beta = 0$ is exactly LTI (Prop. 2)                                              |
| A18     | DC cross-gain of the two forcings (Prop. 3)                                                      |
| A19     | the causality guard bites at the <i>published</i> tuning, and the clairvoyant reference fails it |

disqualify as a finding: our oracle passed the causality guard — assertion A19 re-runs that probe on the published tuning rather than only inside the sweep — which is how we learned the competitor family was mis-specified (Section IV).

We record two calibration facts. Divergent runs are excluded pairwise and reported as counts, because a run that dies early contributes a truncated sum and “wins” computation by dying. And a contrast that moves from  $t = +17.9$  to  $t = +5.4$

when the rival’s tuning grid is densified, or from significant to nothing when a guard is enforced, is a statement about the protocol, not the mechanism; we report both values wherever that happened.

## VIII. CONCLUSIONS

We set out to show that a homeostatic wave field is a good distributed governor for the meta-parameters of MPC, and it is not: in every physically grounded arena it is beaten or matched by a minimal mechanism — and the margin by which we can bound the ties is stated with every claim.

What replaces it is a design lesson with two instances. Where a disturbance travels, deliver its warning with *constant* anticipation: an affine delay line, four parameters, within  $-0.3\%$  to  $+3.7\%$  of the acausal clairvoyant reference, worth  $12.4\%$  against the best local detection, with the proportional family that per-hop designs default to significantly and explicable worse. Where a horizon must be rationed, react to the one event in the suboptimality signal that theory distinguishes — the certificate breaking — through a slow leak: two effective parameters, below the fixed-horizon frontier in a single loop, and up to  $7\%$  ahead of the best fixed horizon in a fleet at equal spent computation under either disturbance-correlation regime — carried by allocation when the loops fail at different times and by timing when they fail together, each mechanism useless or harmful in the other’s regime.

The structural limits say these are not accidents of tuning: nothing built on absolute amplitude can carry information to a meta-parameter, an LTI field offers reshaping rather than information, its natural forcing carries nothing sustained at all, and quantization caps its reach logarithmically. And the five leaks say something about practice that we did not expect to be the paper’s most durable contribution when we began: every one of our arenas was first decided by a defect, the defects were ordinary, the guards are cheap, and two of the five pushed the result *against* the mechanism we were defending. Honest comparisons are not a matter of intention; they are a matter of instrumentation.

## REPRODUCIBILITY

All experiments, figures and the audit suite are generated by archived scripts. A checker, run by the authors before each submission (it is not part of the  $\text{\LaTeX}$  build), regenerates 94 reported quantities from their result files and requires each to occur in the body of this document as a standalone numeric token, anchored on both sides so that a stray 8 does not match inside 0.85. Of the 157 numeric literals in the body with three or more digits or a decimal point, those 94 are regenerated in this sense; 56 more coincide with some value stored in an archived result file, which is a check of existence rather than of provenance; and 7 are declared constants that come from no experiment, such as postal codes and column widths. A mutation test alters, one occurrence at a time, each of the 94 regenerated quantities and re-runs the whole checker; every such mutant is caught. Literals of one or two digits, and the counts in this paragraph, are outside that audit; the checker asserts the counts. Tuning and evaluation seeds are disjoint in every experiment reported.

## ACKNOWLEDGMENT

The authors used Claude (Anthropic) — Claude Opus 4.6 at the start of the project, later Claude Opus 5 and Claude Fable 5.1, as recorded in the repository’s commit history — accessed through the Claude Code command-line assistant, in the preparation of this article. Following the IEEE PSPB Operations Manual, the use is disclosed here in the terms in which it occurred. *Code*. The simulation and analysis code, the plotting scripts, the numerical verifier and its mutation test were written with that assistance, from the protocols, pre-registrations and acceptance criteria specified by the authors, who ran them and are responsible for their outputs. *Text*. All sections were drafted and revised with the same assistance; the authors rewrote and checked every passage and are responsible for every claim it makes. *Adversarial review*. The rounds of adversarial review referred to in the title footnote and in the Introduction were carried out with agents of the same system, directed at this manuscript, at its code and at its archived results; the authors decided which findings to accept and what to do about each. *What is not AI-generated*. No datum reported here was produced, simulated or emulated by an AI system: the reported quantities come from executing the project’s simulation code. No figure contains AI-generated imagery and none was retouched by an AI system; the assistance is confined to the plotting code. No AI system is an author of this article, and the authors take full responsibility for its entire content.

## REFERENCES

- [1] V. Valdivieso, F. G. Montoya, A. Campa-Pinto, and F. M. Arrabal-Campos, “Comment on “multiobjective model predictive control” [Automatica 45 (2009) 2823–2830],” 2026, manuscript, distributed with the code ( [extttbemporad-2009/comment](#)).
- [2] N. Zimmerman, K. Kyprianidis, and C.-F. Lindberg, “Achieving lower district heating network temperatures using feed-forward MPC,” *Materials*, vol. 12, no. 15, p. 2465, 2019.
- [3] E. F. Camacho, F. R. Rubio, M. Berenguel, and L. Valenzuela, “A survey on control schemes for distributed solar collector fields. Part I: Modeling and basic control approaches,” *Solar Energy*, vol. 81, no. 10, pp. 1240–1251, 2007.
- [4] —, “A survey on control schemes for distributed solar collector fields. Part II: Advanced control approaches,” *Solar Energy*, vol. 81, no. 10, pp. 1252–1272, 2007.
- [5] R. Scattolini, “Architectures for distributed and hierarchical model predictive control – a review,” *Journal of Process Control*, vol. 19, no. 5, pp. 723–731, 2009.
- [6] P. D. Christofides, R. Scattolini, D. Muñoz de la Peña, and J. Liu, “Distributed model predictive control: A tutorial review and future research directions,” *Computers & Chemical Engineering*, vol. 51, pp. 21–41, 2013.
- [7] S. Rivero, M. Farina, and G. Ferrari-Trecate, “Plug-and-play decentralized model predictive control for linear systems,” *IEEE Transactions on Automatic Control*, vol. 58, no. 10, pp. 2608–2614, 2013.
- [8] R. R. Negenborn, P.-J. van Overloop, T. Keviczky, and B. De Schutter, “Distributed model predictive control of irrigation canals,” *Networks and Heterogeneous Media*, vol. 4, no. 2, pp. 359–380, 2009.
- [9] P. Barooah, P. G. Mehta, and J. P. Hespanha, “Mistuning-based control design to improve closed-loop stability margin of vehicular platoons,” *IEEE Transactions on Automatic Control*, vol. 54, no. 9, pp. 2100–2113, 2009.
- [10] J. Qi, J. Zhang, and Y. Ding, “Wave equation-based time-varying formation control of multiagent systems,” *IEEE Transactions on Control Systems Technology*, vol. 26, no. 5, pp. 1578–1591, 2018.
- [11] L. Grüne, J. Pannek, M. Seehafer, and K. Worthmann, “Analysis of unconstrained nonlinear MPC schemes with time varying control horizon,” *SIAM Journal on Control and Optimization*, vol. 48, no. 8, pp. 4938–4962, 2010.

- [12] A. J. Krener, "Adaptive horizon model predictive control," *IFAC-PapersOnLine*, vol. 51, no. 13, pp. 31–36, 2018.
- [13] Z. Sun, L. Dai, K. Liu, D. V. Dimarogonas, and Y. Xia, "Robust self-triggered MPC with adaptive prediction horizon for perturbed nonlinear systems," *IEEE Transactions on Automatic Control*, vol. 64, no. 11, pp. 4780–4787, 2019.
- [14] E. G. Gilbert and K. T. Tan, "Linear systems with state and control constraints: The theory and application of maximal output admissible sets," *IEEE Transactions on Automatic Control*, vol. 36, no. 9, pp. 1008–1020, 1991.
- [15] A. Bemporad and D. Muñoz de la Peña, "Multiobjective model predictive control," *Automatica*, vol. 45, no. 12, pp. 2823–2830, 2009.
- [16] S. Boyd and L. Vandenberghe, *Convex Optimization*. Cambridge, UK: Cambridge University Press, 2004.
- [17] L. Grüne and J. Pannek, *Nonlinear Model Predictive Control: Theory and Algorithms*, 2nd ed., ser. Communications and Control Engineering. Cham, Switzerland: Springer, 2017.
- [18] L. Grüne, "Analysis and design of unconstrained nonlinear MPC schemes for finite and infinite dimensional systems," *SIAM Journal on Control and Optimization*, vol. 48, no. 2, pp. 1206–1228, 2009.
- [19] B. Lincoln and A. Rantzer, "Relaxing dynamic programming," *IEEE Transactions on Automatic Control*, vol. 51, no. 8, pp. 1249–1260, 2006.
- [20] S. Holm, "A simple sequentially rejective multiple test procedure," *Scandinavian Journal of Statistics*, vol. 6, no. 2, pp. 65–70, 1979.